



Semaphores in RTOS



Introduction

- ▶ Multiple concurrent threads (***tasks***) of execution within an application must be able to
 - ▶ ***synchronize their execution*** and
 - ▶ ***coordinate mutually exclusive access to shared resources.***
- ▶ This can be accomplished using RTOS kernels ***semaphore object*** and ***associated semaphore management*** services.



Defining Semaphores

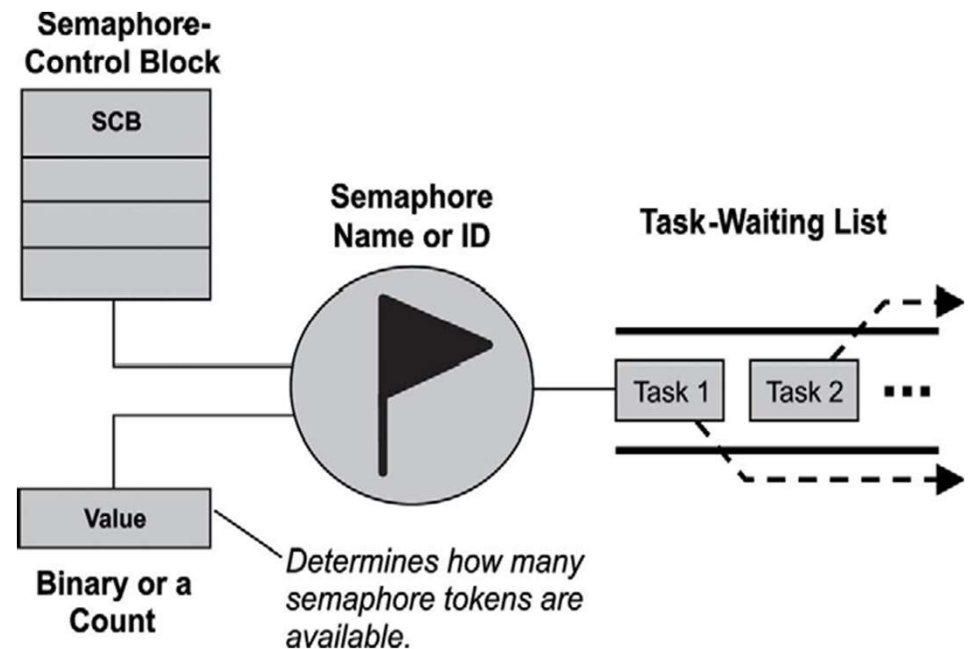
- ▶ Definition

- ▶ A *semaphore* (sometimes called a *semaphore token*) is a kernel object that one or more threads of execution can acquire or release for the purposes of synchronization or mutual exclusion.



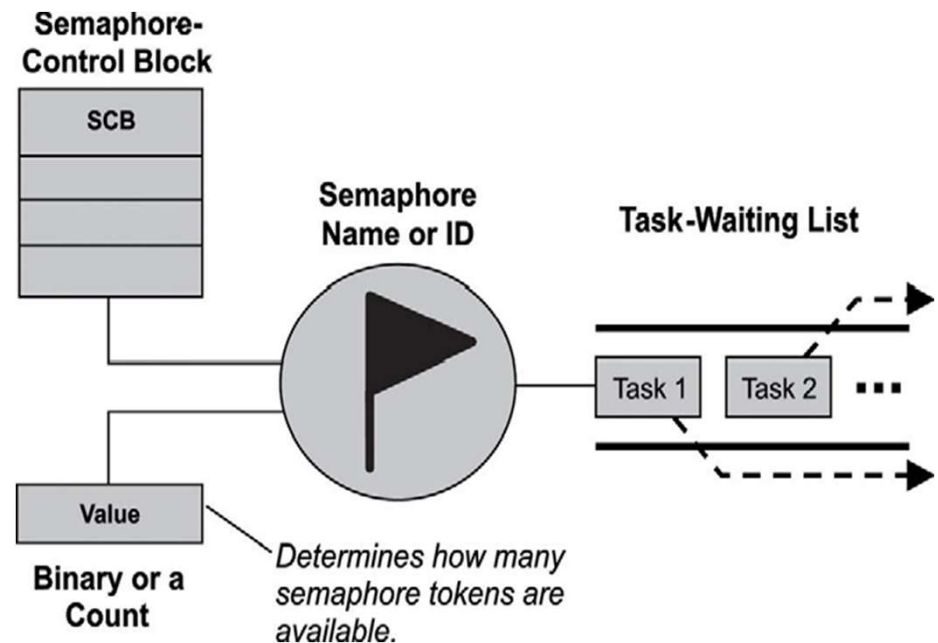
Defining Semaphores

- ▶ When a semaphore is first created, the kernel assigns to it an associated
 - ▶ semaphore control block (SCB),
 - ▶ a unique ID,
 - ▶ a value (binary or a count), and
 - ▶ a task-waiting list



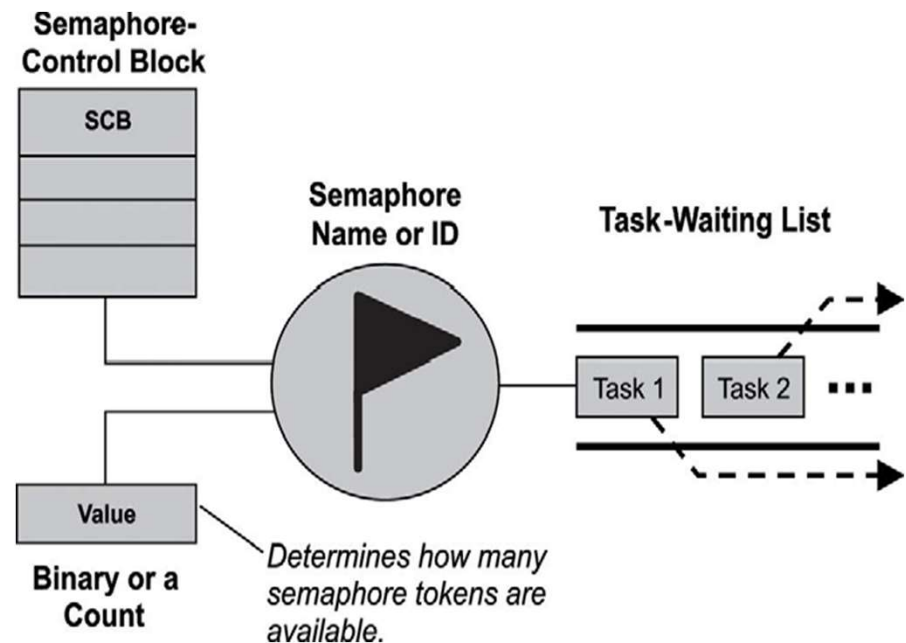
Defining Semaphores

- ▶ Counting semaphore
 - ▶ A counter is maintained for token counting
 - ▶ When a task acquires the semaphore, the token count is decremented; as a task releases the semaphore, the count is incremented.



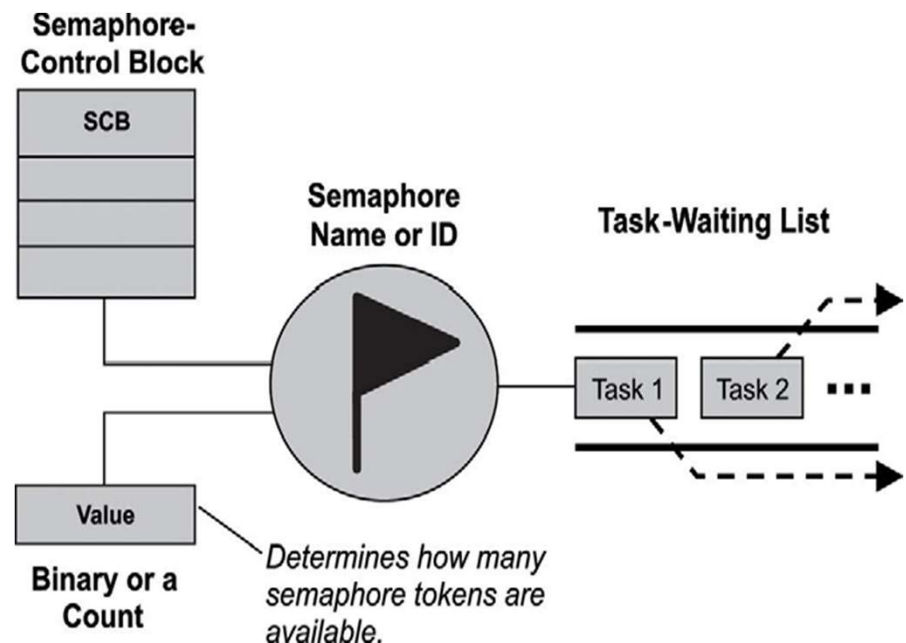
Defining Semaphores

- ▶ If the token count reaches 0 that no more resources is available and the requesting task is ***blocked***.



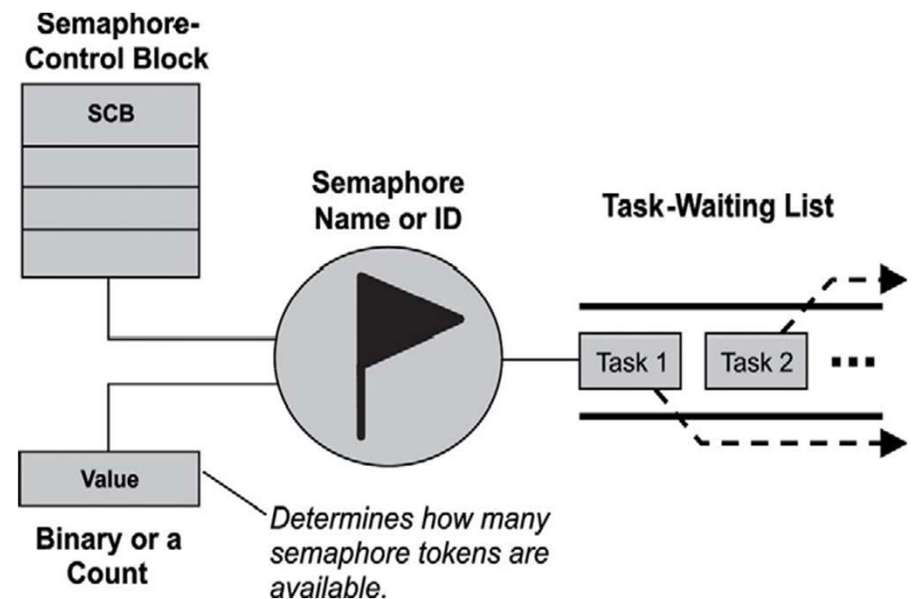
Defining Semaphores

- ▶ The ***task-waiting list*** tracks all blocked tasks of the semaphore.
- ▶ The list is implemented either in FIFO order or highest priority first order.



Defining Semaphores

- ▶ When an semaphore/resource becomes available, the kernel allows the first task in the task- waiting list to acquire it.
 - ▶ if the task has the highest priority, the task moves to the running state, otherwise it moves
 - ▶ to the ready state wait for its turn
 - ▶ Note that the exact implementation of a task-waiting list can vary from one kernel to another.




Defining Semaphores

- ▶ A kernel can support many different types of semaphores, including
 - ▶ binary,
 - ▶ counting, and
 - ▶ mutual-exclusion (mutex) semaphores.




Binary Semaphores

- ▶ A *binary semaphore* can have a value of either 0 or 1
 - ▶ 0, semaphore is *unavailable (or empty)*
 - ▶ 1, the semaphore is *available (or full)*.



Binary Semaphores

- ▶ Binary semaphores are treated as **global resources**,
 - ▶ they are shared among **all tasks** that need them.
 - ▶ Making the semaphore a global resource allows any task to release it, even if the task did not initially acquire it???



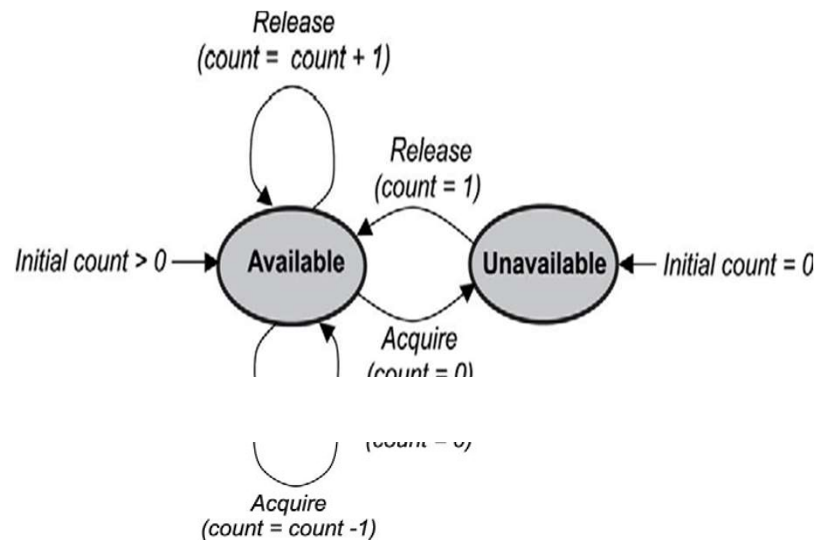
Counting Semaphores

- ▶ A *counting semaphore* uses a count to allow it to be acquired or released multiple times.




Counting Semaphores

- ▶ If the initial count is
 - ▶ 0, the counting semaphore is created in the unavailable state.
 - ▶ greater than 0, the semaphore is created in the available state, and the number of tokens it has equals its count



Counting Semaphores

- ▶ counting semaphores are **global resources** that can be shared by all tasks that need them




Mutual Exclusion (Mutex) Semaphores

- ▶ A *mutual exclusion (mutex) semaphore* is a special binary semaphore that supports
 - ▶ **ownership**,
 - ▶ recursive access,
 - ▶ task deletion safety, and
 - ▶ one or more protocols for avoiding problems inherent to mutual exclusion.




Mutual Exclusion (Mutex) Semaphores

- ▶ the states of a mutex are
 - ▶ 0 *unlocked* or
 - ▶ 1 *locked*



Mutual Exclusion (Mutex) Semaphores

- ▶ A mutex is initially created in the ***unlocked state***, in which it can be acquired by a task.
- ▶ After being acquired, the mutex moves to the ***locked state***



Mutual Exclusion (Mutex) Semaphores


- ▶ **Mutex Ownership**

- ▶ *Ownership* of a mutex is gained when a **task** first locks the mutex by acquiring it.
- ▶ Conversely, a task loses ownership of the mutex when it unlocks it by releasing it.




Mutual Exclusion (Mutex) Semaphores

- ▶ Recursive Locking
 - ▶ allows the task that owns the mutex to acquire it **multiple times** in the locked state
 - ▶ The mutex with recursive locking is called a **recursive mutex**



Mutual Exclusion (Mutex) Semaphores

- ▶ Recursive Locking
 - ▶ A recursive mutex allows *nested attempts to lock the mutex to succeed*, rather than cause *deadlock*,



See 6.4.5 Recursive Shared-Resource-Access Synchronization



Mutual Exclusion (Mutex) Semaphores

- ▶ Difference between Counting Semaphore and Mutex
 - ▶ The count used for the mutex tracks the number of times that ***the task owning the mutex*** has locked or unlocked the mutex.
 - ▶ The count used for the counting semaphore tracks the number of tokens that have been ***acquired or released by any task***.



Mutual Exclusion (Mutex) Semaphores

- ▶ Task Deletion Safety
 - ▶ Premature task deletion is avoided by using *task deletion locks* when a task locks and unlocks a mutex
 - ▶ Enabling this capability within a mutex ensures that *while a task owns the mutex, the task cannot be deleted*



Mutual Exclusion (Mutex) Semaphores

- ▶ Priority Inversion Avoidance
 - ▶ Priority inversion occurs when a higher priority task is blocked and is waiting for a resource being used by a lower priority task, which has itself been preempted by an unrelated medium-priority task. In this situation, the higher priority task's priority level has effectively been inverted to the lower priority task's level.



Mutual Exclusion (Mutex) Semaphores

- ▶ Priority Inversion Avoidance
- ▶ Two common protocols used for avoiding priority inversion include:
 - ▶ **priority inheritance protocol**
 - ▶ **ceiling priority protocol**

apply to the task that owns the mutex.



Mutual Exclusion (Mutex) Semaphores

- ▶ **priority inheritance protocol—**
 - ▶ ensures that the priority level of the lower priority task that has acquired the mutex is ***raised to that of the higher priority task that has requested the mutex*** when inversion happens.
 - ▶ The priority of the raised task is lowered to its original value after the task releases the mutex that the higher priority task requires.



Mutual Exclusion (Mutex) Semaphores

- ▶ **ceiling priority protocol—**
 - ▶ ensures that the priority level of the task that acquires the mutex is automatically *set to the highest priority of all possible tasks* that might request that mutex when it is first acquired until it is released.



Typical Semaphore Operations

- ▶ Typical operations that developers might want to perform with the semaphores in an application include:
 - ▶ creating and deleting semaphores,
 - ▶ acquiring and releasing semaphores,
 - ▶ clearing a semaphore's task-waiting list, and
 - ▶ getting semaphore information.



Creating and Deleting Semaphores

► Create and Delteting

Table 6.1: Semaphore creation and deletion operations.

Operation	Description
Create	Creates a semaphore
Delete	Deletes a semaphore

```
create_binary_semaphore(initial_state, task_waiting_order);  
create_counting_semaphore(initial_count, task_waiting_order);  
create_mutex_semaphore(task_waiting_order);
```



Creating and Deleting Semaphores

- ▶ different calls might be used for creating binary, counting, and mutex semaphores, as follows:
 - ▶ **binary**—specify the initial semaphore state and the task-waiting order.
 - ▶ **counting**—specify the initial semaphore count and the task-waiting order.

```
delete_semaphore(semaphore_id);
```



Creating and Deleting Semaphores

- ▶ different calls might be used for creating binary, counting, and mutex semaphores, as follows:
 - ▶ **mutex**—specify the task-waiting order and enable task deletion safety, recursion, and priority-inversion avoidance protocols, if supported.

```
create_mutex_semaphore(task_waiting_order, enable_task_deletion_safety,  
enable_recursion, enable_priority_inversion_avoidance);
```



Creating and Deleting Semaphores

- ▶ Deleting (implementation dependent)
 - ▶ Semaphores can be deleted from within any task by specifying their IDs and making semaphore-deletion calls.
 - ▶ Deleting a semaphore is not the same as releasing it. When a semaphore is deleted, blocked tasks in its task-waiting list are unblocked and moved either to the ready state or to the running state (if the unblocked task has the highest priority).
 - ▶ Any tasks (not in the task waiting list), however, that try to acquire the deleted semaphore return with an error because the semaphore no longer exists.



Creating and Deleting Semaphores

- ▶ Deleting

- ▶ Additionally, do not delete a semaphore while it is in use (e.g., acquired). This action might result in data corruption or other serious problems if the semaphore is protecting a shared resource or a critical section of code.



Acquiring and Releasing Semaphores

Table 6.2: Semaphore acquire and release operations.

Operation	Description
Acquire (<i>take/ sm_p/ pend/ lock</i>)	Acquire a semaphore token
Release (<i>give/ sm_v/ post/ unlock</i>)	Release a semaphore token



Acquiring and Releasing Semaphores

- ▶ Tasks typically make a request to **acquire** a semaphore in one of the following ways:
 - ▶ **Wait forever**—task remains blocked until it is able to acquire a semaphore.
 - ▶ **Wait with a timeout**—task remains blocked until it is able to acquire a semaphore or until a set interval of time, called the *timeout interval* , passes.
 - ▶ **Do not wait**—task makes a request to acquire a semaphore token, but, if one is not available, the task does not block.



Acquiring and Releasing Semaphores

- ▶ Release

- ▶ Any task can ***release*** a binary or counting semaphore;
- ▶ however, a mutex can only be released (unlocked) by the task that first acquired (locked) it.



Acquiring and Releasing Semaphores

- ▶ Release

- ▶ Note that incorrectly releasing a binary or counting semaphore can result in losing mutually exclusive access to a shared resource or in an I/O device malfunction.



Acquiring and Releasing Semaphores

- ▶ ISR and Semaphores
 - ▶ Note that ISRs can also release binary and counting semaphores.
 - ▶ Note that most kernels ***do not*** support ISRs locking and unlocking mutexes, as it is not meaningful to do so from an ISR.
 - ▶ It is also not meaningful to acquire either binary or counting semaphores inside an ISR.



Clearing Semaphore Task-Waiting Lists

- ▶ To ***clear*** all tasks waiting on a semaphore task-waiting list

Table 6.3: Semaphore unblock operations.

Operation	Description
Flush	Unblocks all tasks waiting on a semaphore



Clearing Semaphore Task-Waiting Lists

-
- ▶ ***thread rendezvous*** that implement *flush* in synchronization tasks
 - ▶ multiple tasks to complete certain activities first and then
 - ▶ block while trying to acquire a common semaphore that is made unavailable, after
 - ▶ the last task finishes doing what it needs to, the task can execute a semaphore flush operation on the common semaphore.
 - ▶ This operation frees all tasks waiting in the semaphore's task waiting list.



Getting Semaphore Information

- ▶ obtain semaphore information to perform monitoring or debugging

Table 6.4: Semaphore information operations.

Operation	Description
Show info	Show general information about semaphore
Show blocked tasks	Get a list of IDs of tasks that are blocked on a semaphore



Typical Semaphore Use

- ▶ Semaphores are useful either for synchronizing execution of multiple tasks or for coordinating access to a shared resource.





Typical Semaphore Use

- ▶ Semaphores are useful
 - ▶ wait-and-signal synchronization,
 - ▶ multiple-task wait-and-signal synchronization,
 - ▶ credit-tracking synchronization,
 - ▶ single shared-resource-access synchronization,
 - ▶ recursive shared-resource-access synchronization, and
 - ▶ multiple shared-resource-access synchronization.
- and etc.




Wait-and-Signal Synchronization

- ▶ Two tasks can communicate for the purpose of synchronization without exchanging data.



Wait-and-Signal Synchronization

1. binary semaphore is initially unavailable (value of 0).
2. tWaitTask has higher priority and runs first and blocked by the semaphore
3. lower priority tSignalTask a chance to run and release semaphore
4. tWaitTask unblocked and preempts tSignalTask and starts to execute




```
tWaitTask ( ) {  
    :  
    Acquire binary semaphore token  
    :  
}  
tSignalTask ( ) {  
    .  
    Release binary semaphore token  
    :  
}
```




Multiple-Task Wait-and-Signal Synchronization

- ▶ To coordinate the synchronization of more than two tasks, use the flush operation on the task-waiting list of a binary semaphore,



Multiple-Task Wait-and-Signal Synchronization

1. the binary semaphore is initially unavailable (value of 0).
2. The higher priority **tWaitTasks 1, 2, and 3** execute, acquire the unavailable semaphore, and blocked.
3. **tSignalTask** execute and flush the semaphore
4. one of the higher priority tWaitTasks preempts tSignalTask and starts to execute



```

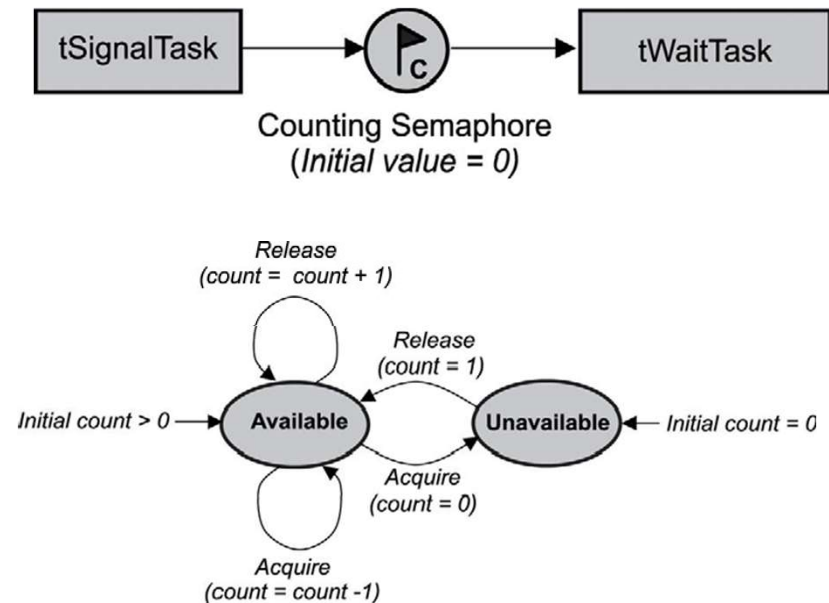
tWaitTask () {
    Acquire binary semaphore token
    :
}

tSignalTask () {
    :
    Flush binary semaphore's task-waiting list
    :
}
  
```




Credit-Tracking Synchronization

1. counting semaphore's count is initially 0, making it unavailable
2. lower priority *tWaitTask* tries to **acquire** this semaphore but blocks until *tSignalTask* makes the semaphore available by performing a **release**



Credit-Tracking Synchronization

3. Even then, **tWaitTask** will wait in the ready state until the higher priority **tSignalTask** eventually relinquishes the CPU by making a blocking call or delaying itself,
4. Because **tSignalTask** is set to a higher priority and executes at its own rate, it might increment the counting semaphore multiple times before **tWaitTask** starts processing the first request.




```
tWaitTask () {  
    :  
    Acquire counting semaphore token  
    :  
}  
  
tSignalTask () {  
    :  
    Release counting semaphore token  
    :  
}
```



Credit-Tracking Synchronization

5. Hence, the counting semaphore allows a credit buildup of the number of times that the **tWaitTask** can execute before the semaphore becomes unavailable




tSignalTask is the source.
releasing token
tWaitTask is the sink,
Acquire (consume) the token




Credit-Tracking Synchronization

6. when **tSignalTask's** rate of releasing the semaphore tokens slows, tWaitTask can catch up and eventually deplete the count until the counting semaphore is empty




Credit-Tracking Synchronization

7. At this point, **tWaitTask** blocks again at the counting semaphore, waiting for **tSignalTask** to release the semaphore again




Credit-Tracking Synchronization

8. useful if **tSignalTask** releases semaphores in *bursts*, giving **tWaitTask** the chance to catch up every once in a while



Credit-Tracking Synchronization

- ▶ ISR and DST
 - ▶ ISR that acts in a similar way to the signaling task
 - ▶ DSR/Wait task with lower priority offloads works for ISR




Single Shared-Resource-Access Synchronization

- ▶ To provide for mutually exclusive access to a shared resource, that a
 - ▶ memory location,
 - ▶ a data structure, or an
 - ▶ I/O device
 - ▶ essentially anything that might have to be shared between two or more concurrent threads of execution.
- ▶ A semaphore can be used to serialize access to a shared resource,




Single Shared-Resource-Access Synchronization

1. a binary semaphore is initially created in the available state (value = 1) and is used to protect the shared resource



Single Shared-Resource-Access Synchronization

2. To access the shared resource, task 1 or 2 needs to first successfully acquire the binary semaphore before reading from or writing to the shared resource



```
tAccessTask () {  
    :  
    Acquire binary semaphore token  
    Read or write to shared resource  
    Release binary semaphore token  
    :  
} # task1 and task 2
```




Single Shared-Resource-Access Synchronization

- ❑ **Pitfall!!**

- ❑ A third party task can accidentally release the binary semaphore, even one that never acquired the semaphore in the first place.

- ❑ If this is the case, both **tAccessTask 1** and **tAccessTask 2** could end up acquiring the semaphore and accessing the resource at the same time

- ❑ Using **mutex** to avoid




```
tAccessTask () {  
    :  
    Acquire binary semaphore token  
    Read or write to shared resource  
    Release binary semaphore token  
    :  
} # task1 and task 2
```



Single Shared-Resource-Access Synchronization

▮ MUTEX

- ▮ Mutex is a special semaphore that only the task acquired it can release it.

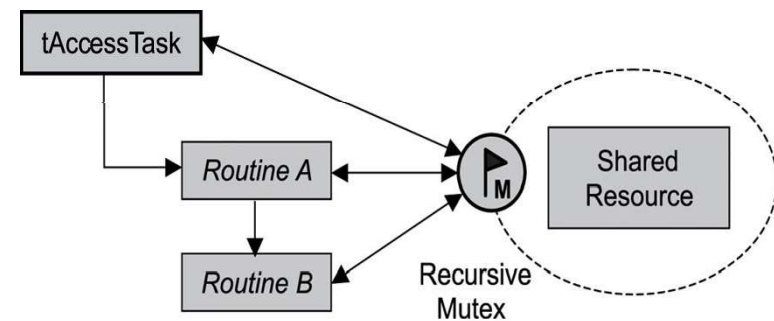


```
tAccessTask () {  
    :  
    Acquire binary semaphore token  
    Read or write to shared resource  
    Release binary semaphore token  
    :  
} # task1 and task 2
```



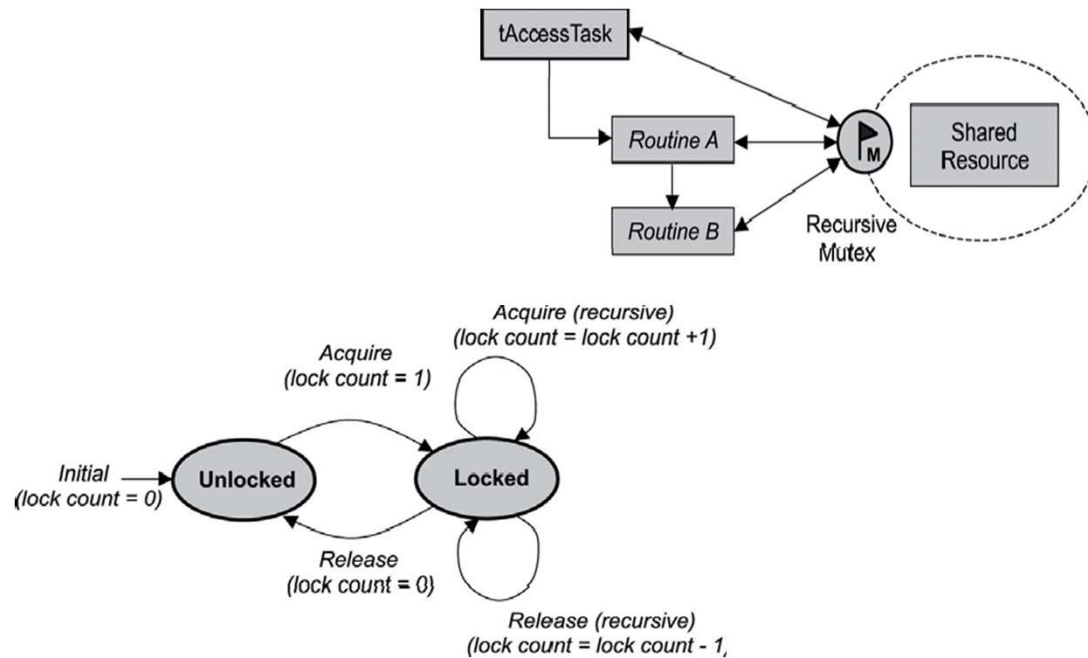
Recursive Shared-Resource-Access Synchronization

- ▶ Task to access a shared resource recursively.
 - ▶ For example, **tAccessTask** calls Routine A that calls Routine B, and all three need access to the same shared resource,
 - ▶ tasks would end up blocking, causing a deadlock



Recursive Shared-Resource-Access Synchronization

```
tAccessTask () {  
  :  
  Acquire mutex  
  Access shared resource  
  Call Routine A  
  Release mutex  
  :  
}  
  
Routine A () {  
  :  
  Acquire mutex  
  Access shared resource  
  Call Routine B  
  Release mutex  
  :  
}  
  
Routine B () {  
  :  
  Acquire mutex  
  Access shared resource  
  Release mutex  
  :  
}
```




recursive mutex. After `tAccessTask` locks the mutex, the task owns it. Additional attempts from the task itself or from routines that it calls to lock the mutex succeed.



Multiple Shared-Resource-Access Synchronization



- ▶ When multiple equivalent shared resources are used, a counting semaphore comes in handy
- ▶ The counting semaphore's count is initially set to the number of equivalent shared resources



Multiple Shared-Resource-Access Synchronization

```
tAccessTask () {  
    :  
    Acquire a counting semaphore token  
    Read or Write to shared resource  
    Release a counting semaphore token  
    :  
}
```

similar code is used for
tAccessTask 1, 2, and 3



Multiple Shared-Resource-Access Synchronization

```
tAccessTask () {
```

```
:
```

```
    Acquire first mutex in non-blocking way
```

```
        If not successful then acquire 2nd mutex in a blocking way
```

```
    Read or Write to shared resource
```

```
    Release the acquired mutex
```

```
:
```

```
}
```

- ▶ similar code is used for tAccessTask 1, 2, and 3
- ▶ Pseudo code for multiple tasks accessing equivalent shared resources using mutexes



Đoạn mã giả định cho việc nhiều tác vụ (tAccessTask1, tAccessTask2, tAccessTask3, vv.)

```
tAccessTask1():  
    while True:  
        # Attempt to acquire the first mutex in a non-blocking way  
        if mutex1.try_acquire():  
            # First mutex acquired successfully  
            break # Exit the loop  
        else:  
            # First mutex not acquired, try to acquire the second mutex in a blocking way  
            mutex2.acquire()  
  
    # Access the shared resource (read or write)  
    # [Code to read from or write to the shared resource]  
  
    # Release the acquired mutex  
    mutex1.release() # Release the first mutex  
    mutex2.release() # Release the second mutex
```


Đoạn mã giả định cho việc nhiều tác vụ (tAccessTask1, tAccessTask2, tAccessTask3, vv.)_

```
tAccessTask2():  
    while True:  
        # Attempt to acquire the first mutex in a non-blocking way  
        if mutex1.try_acquire():  
            # First mutex acquired successfully  
            break # Exit the loop  
        else:  
            # First mutex not acquired, try to acquire the second mutex in a blocking  
way            mutex2.acquire()  
  
        # Access the shared resource (read or write)  
        # [Code to read from or write to the shared resource]  
  
        # Release the acquired mutex  
        mutex1.release() # Release the first mutex  
        mutex2.release() # Release the second mutex
```

Đoạn mã giả định cho việc nhiều tác vụ (tAccessTask1, tAccessTask2, tAccessTask3, vv.)_

```
tAccessTask3():
    while True:
        # Attempt to acquire the first mutex in a non-blocking way
        if mutex1.try_acquire():
            # First mutex acquired successfully
            break # Exit the loop
        else:
            # First mutex not acquired, try to acquire the second mutex in a blocking way
            mutex2.acquire()

    # Access the shared resource (read or write)
    # [Code to read from or write to the shared resource]

    # Release the acquired mutex
    mutex1.release() # Release the first mutex
    mutex2.release() # Release the second mutex
```